




Module 7: Storage

AWS Academy Cloud Foundations

Core AWS services



Amazon Virtual Private Cloud (Amazon VPC)



Amazon Elastic Compute Cloud (Amazon EC2)



Storage



AWS Identity and Access Management (IAM)




Amazon Relational Database Service



Amazon DynamoDB

Database

Storage




Amazon Elastic Block Store
(Amazon EBS)

AWS storage options: Block storage versus object storage




What if you want to change **one character** in a 1-GB file?



Block storage

Change one block (piece of the file)
that contains the character



Object storage

Entire file must be updated

Amazon EBS

Amazon EBS enables you to **create individual storage volumes** and **attach them** to an Amazon EC2 instance:

- Amazon EBS offers block-level storage.
- Volumes are automatically replicated within its Availability Zone.
- It can be backed up automatically to Amazon S3 through snapshots.
- Uses include –
 - Boot volumes and storage for Amazon Elastic Compute Cloud (Amazon EC2) instances
 - Data storage with a file system
 - Database hosts
 - Enterprise applications

Amazon EBS volume types

Maximum Volume Size
Maximum IOPS/Volume
Maximum
Throughput/Volume


Solid State Drives (SSD)		Hard Disk Drives (HDD)	
General Purpose	Provisioned IOPS	Throughput-Optimized	Cold
16 TiB	16 TiB	16 TiB	16 TiB
16,000	64,000	500	250
250 MiB/s	1,000 MiB/s	500 MiB/s	250 MiB/s

Amazon EBS volume type use cases

Solid State Drives (SSD)		Hard Disk Drives (HDD)	
General Purpose	Provisioned IOPS	Throughput-Optimized	Cold
<ul style="list-style-type: none"> This type is recommended for most workloads 	<ul style="list-style-type: none"> Critical business applications that require sustained IOPS performance, or more than 16,000 IOPS or 250 MiB/second of throughput per volume 	<ul style="list-style-type: none"> Streaming workloads that require consistent, fast throughput at a low price 	<ul style="list-style-type: none"> Throughput-oriented storage for large volumes of data that is infrequently accessed
<ul style="list-style-type: none"> System boot volumes 	<ul style="list-style-type: none"> Large database workloads 	<ul style="list-style-type: none"> Big data 	<ul style="list-style-type: none"> Scenarios where the lowest storage cost is important
<ul style="list-style-type: none"> Virtual desktops 		<ul style="list-style-type: none"> Data warehouses 	<ul style="list-style-type: none"> It cannot be a boot volume
<ul style="list-style-type: none"> Low-latency interactive applications 		<ul style="list-style-type: none"> Log processing 	
<ul style="list-style-type: none"> Development and test environments 		<ul style="list-style-type: none"> It cannot be a boot volume 	

Amazon EBS features

- Snapshots –
 - Point-in-time snapshots
 - Recreate a new volume at any time
- Encryption –
 - Encrypted Amazon EBS volumes
 - No additional cost
- Elasticity –
 - Increase capacity
 - Change to different types



Storage



Amazon Simple Storage Service (Amazon S3)

Amazon S3 overview



- Data is stored as objects in buckets
- Virtually unlimited storage
 - Single object is limited to 5 TB
- Designed for 11 9s of durability
- Granular access to bucket and objects




Amazon S3 storage classes

Amazon S3 offers a range of object-level storage classes that are designed for different use cases:

- Amazon S3 Standard
- Amazon S3 Intelligent-Tiering
- Amazon S3 Standard-Infrequent Access (Amazon S3 Standard-IA)
- Amazon S3 One Zone-Infrequent Access (Amazon S3 One Zone-IA)
- Amazon S3 Glacier
- Amazon S3 Glacier Deep Archive

Amazon S3 bucket URLs (two styles)



To upload your data:

1. Create a **bucket** in an AWS Region.
2. Upload almost any number of **objects** to the bucket.

Bucket path-style URL endpoint:

<https://s3.ap-northeast-1.amazonaws.com/bucket-name>

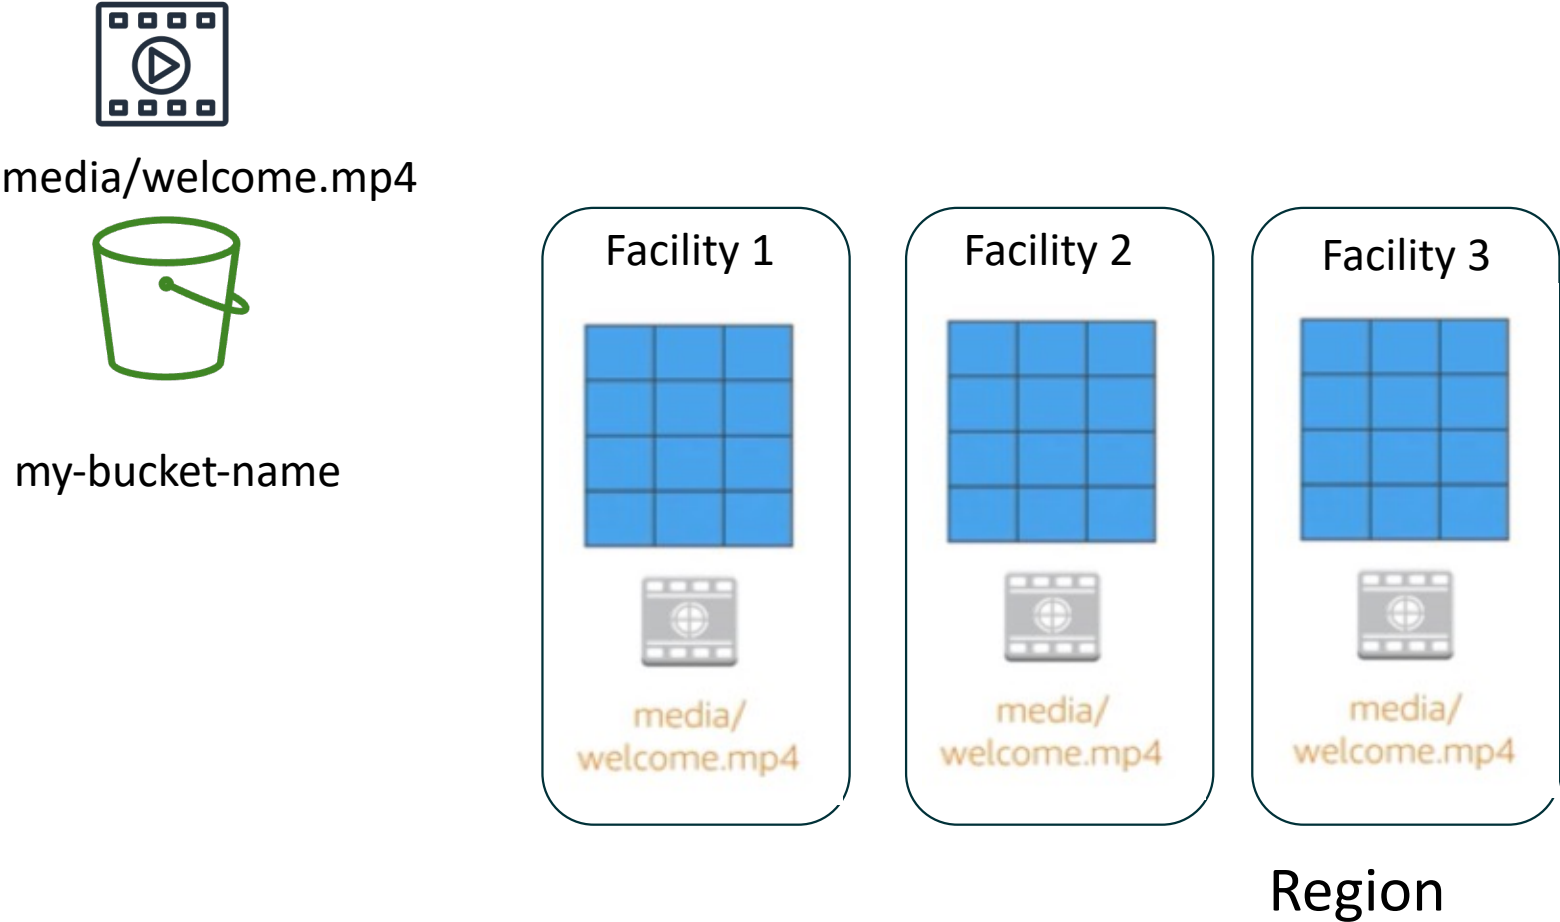
Region code Bucket name

Bucket virtual hosted-style URL endpoint:


<https://bucket-name.s3-ap-northeast-1.amazonaws.com>

Bucket name Region code

Data is redundantly stored in the Region



Access the data anywhere



AWS Management
Console



AWS Command Line
Interface



SDK

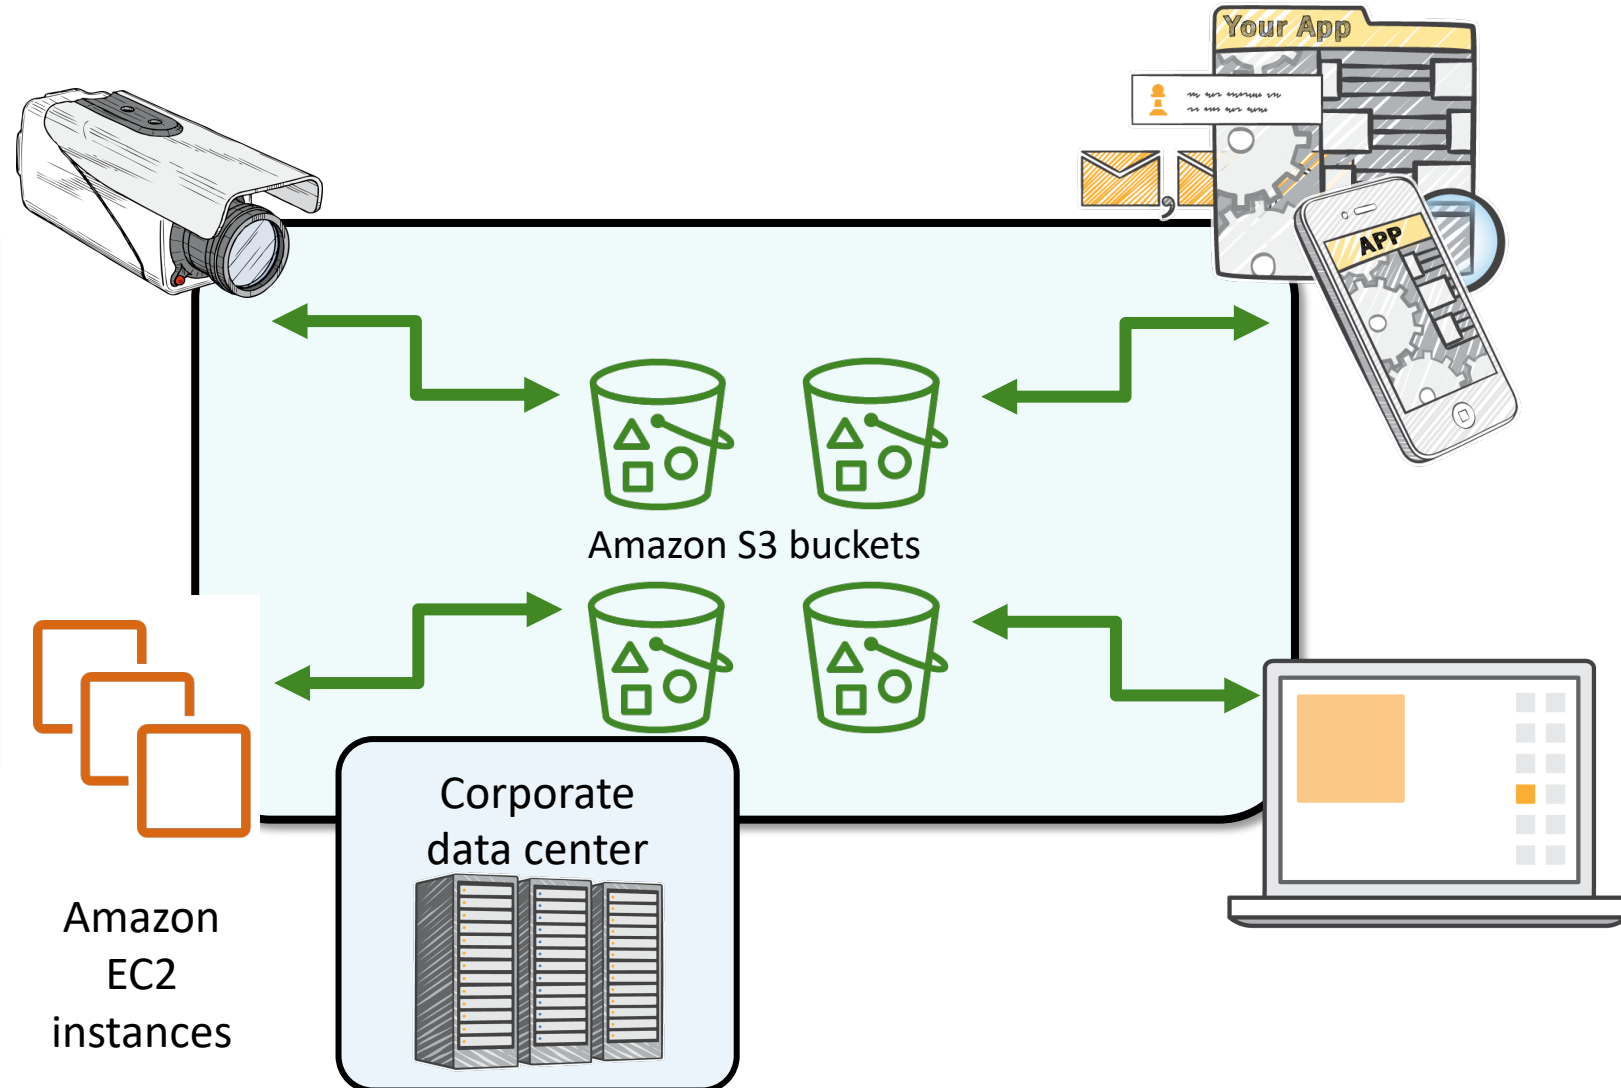
Common use cases

- Storing application assets
- Static web hosting
- Backup and disaster recovery (DR)
- Staging area for big data
- *Many more....*



Amazon S3 common scenarios

- Backup and storage
- Application hosting
- Media hosting
- Software delivery



Amazon S3: Storage pricing (1 of 2)

To estimate Amazon S3 costs, consider the following:

1. Storage class type –

- Standard storage is designed for:
 - 11 9s of durability
 - Four 9s of availability
- S3 Standard-Infrequent Access (S-IA) is designed for:
 - 11 9s of durability
 - Three 9s of availability

2. Amount of storage –

- The number and size of objects

Amazon S3: Storage pricing (2 of 2)

3. Requests –

- The number and type of requests (**GET, PUT, COPY**)
- Type of requests:
 - Different rates for GET requests than other requests.

4. Data transfer –

- Pricing is based on the amount of data that is transferred out of the Amazon S3 Region
 - Data transfer in is free, but you incur charges for data that is transferred out.

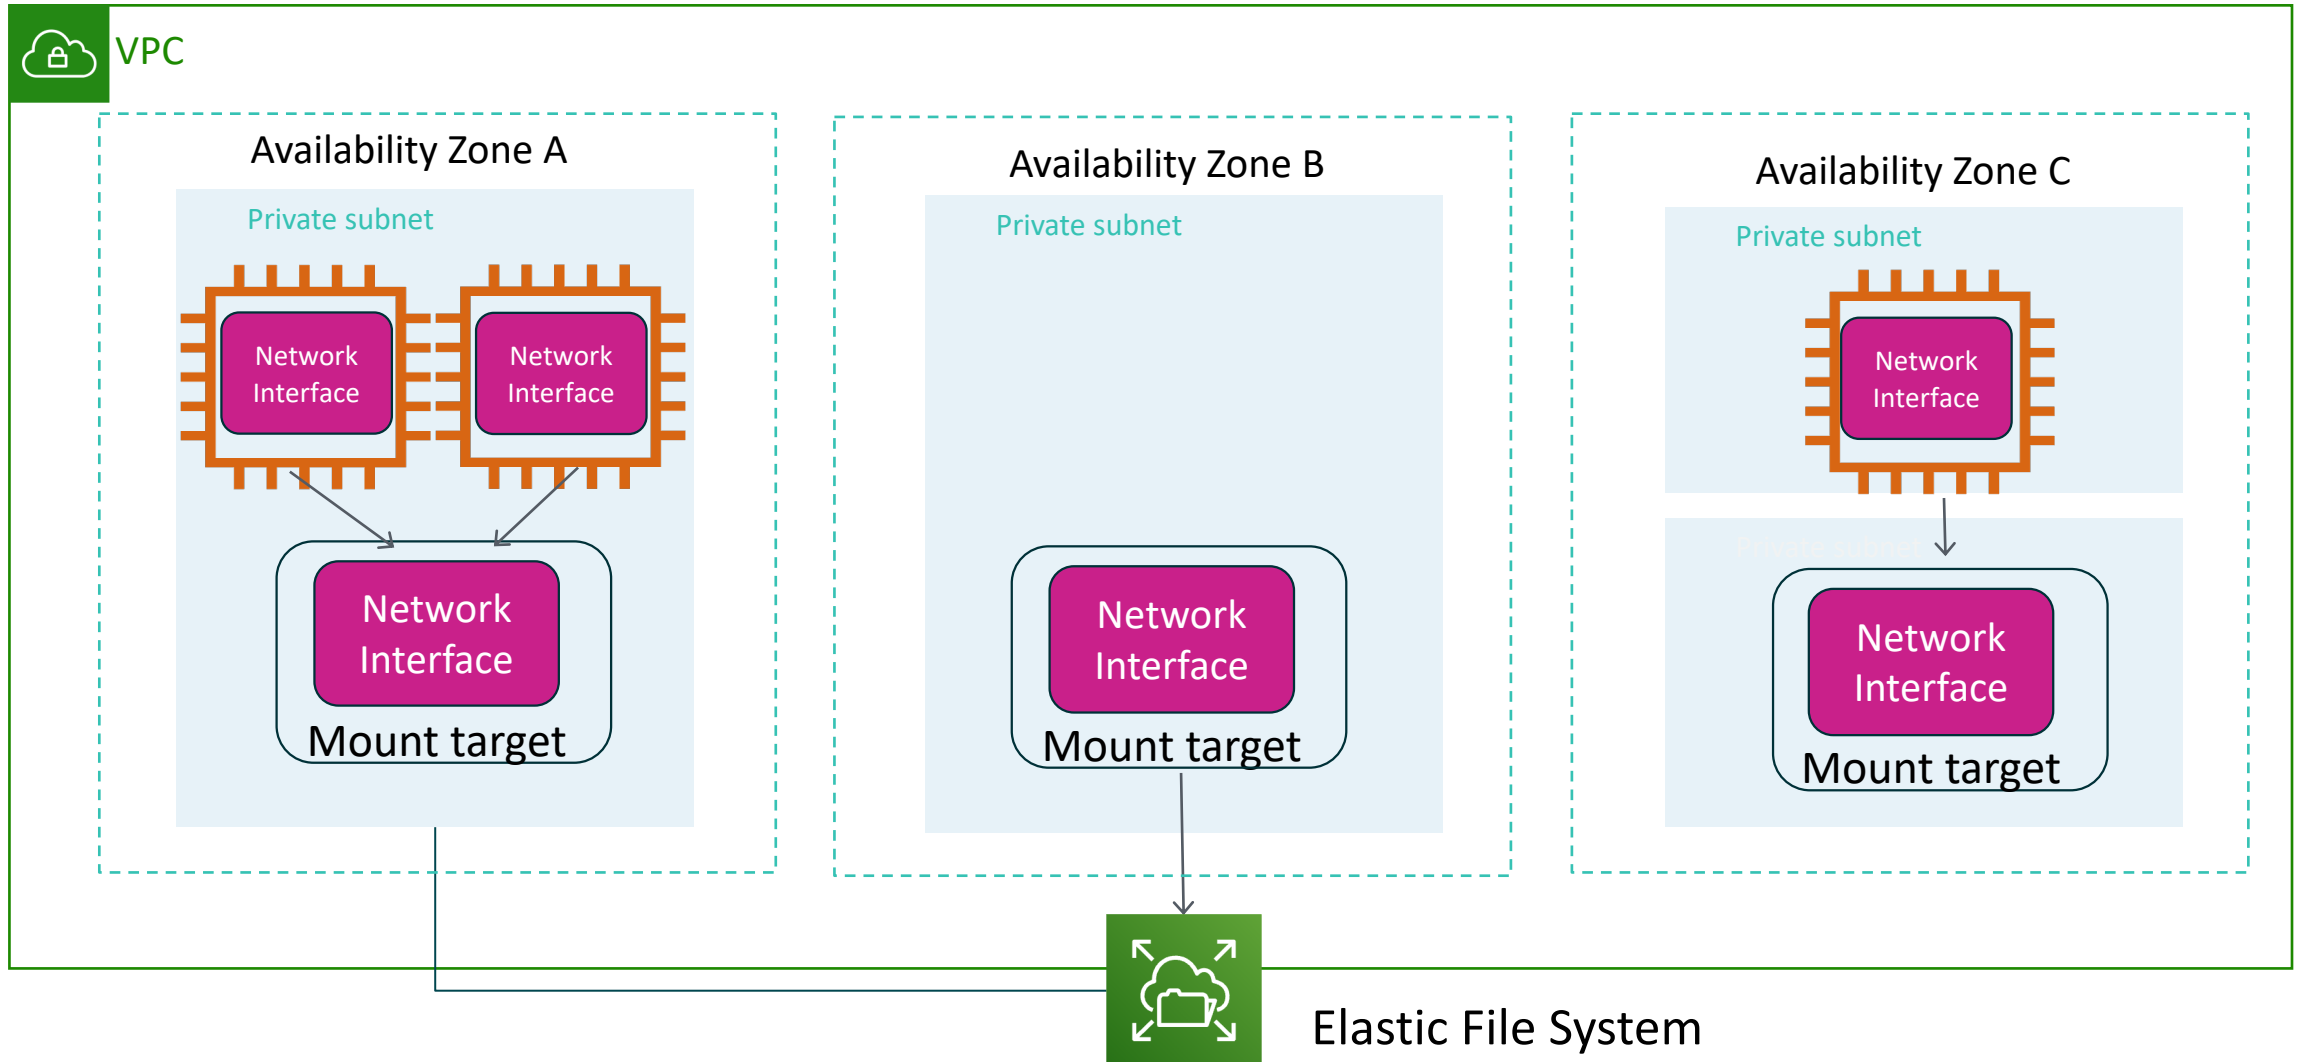


Amazon Elastic File System (Amazon EFS)

Amazon EFS features

- File storage in the AWS Cloud
- Works well for big data and analytics, media processing workflows, content management, web serving, and home directories
- Petabyte-scale, low-latency file system
- Shared storage
- Elastic capacity
- Supports Network File System (NFS) versions 4.0 and 4.1 (NFSv4)
- Compatible with all Linux-based AMIs for Amazon EC2

Amazon EFS architecture



Storage



Amazon S3 Glacier


Amazon S3 Glacier review

Amazon S3 Glacier is a **data archiving service** that is designed for **security**, **durability**, and an **extremely low cost**.

- Amazon S3 Glacier is designed to provide 11 9s of durability for objects.
- It supports the encryption of data in transit and at rest through Secure Sockets Layer (SSL) or Transport Layer Security (TLS).
- The Vault Lock feature enforces compliance through a policy.
- Extremely low-cost design works well for long-term archiving.
 - Provides three options for access to archives—expedited, standard, and bulk—retrieval times range from a few minutes to several hours.

Amazon S3 Glacier

- Storage service for low-cost data archiving and long-term backup
- You can configure lifecycle archiving of Amazon S3 content to Amazon S3 Glacier
- Retrieval options –
 - Standard: 3–5 hours
 - Bulk: 5–12 hours
 - Expedited: 1–5 minutes



Amazon S3 Glacier use cases



Media asset archiving



Healthcare information archiving



Regulatory and compliance archiving



Scientific data archiving




Digital preservation



Magnetic tape replacement

Using Amazon S3 Glacier




RESTful
web services

Java or .NET
SDKs

Amazon S3 with
lifecycle policies

Lifecycle policies

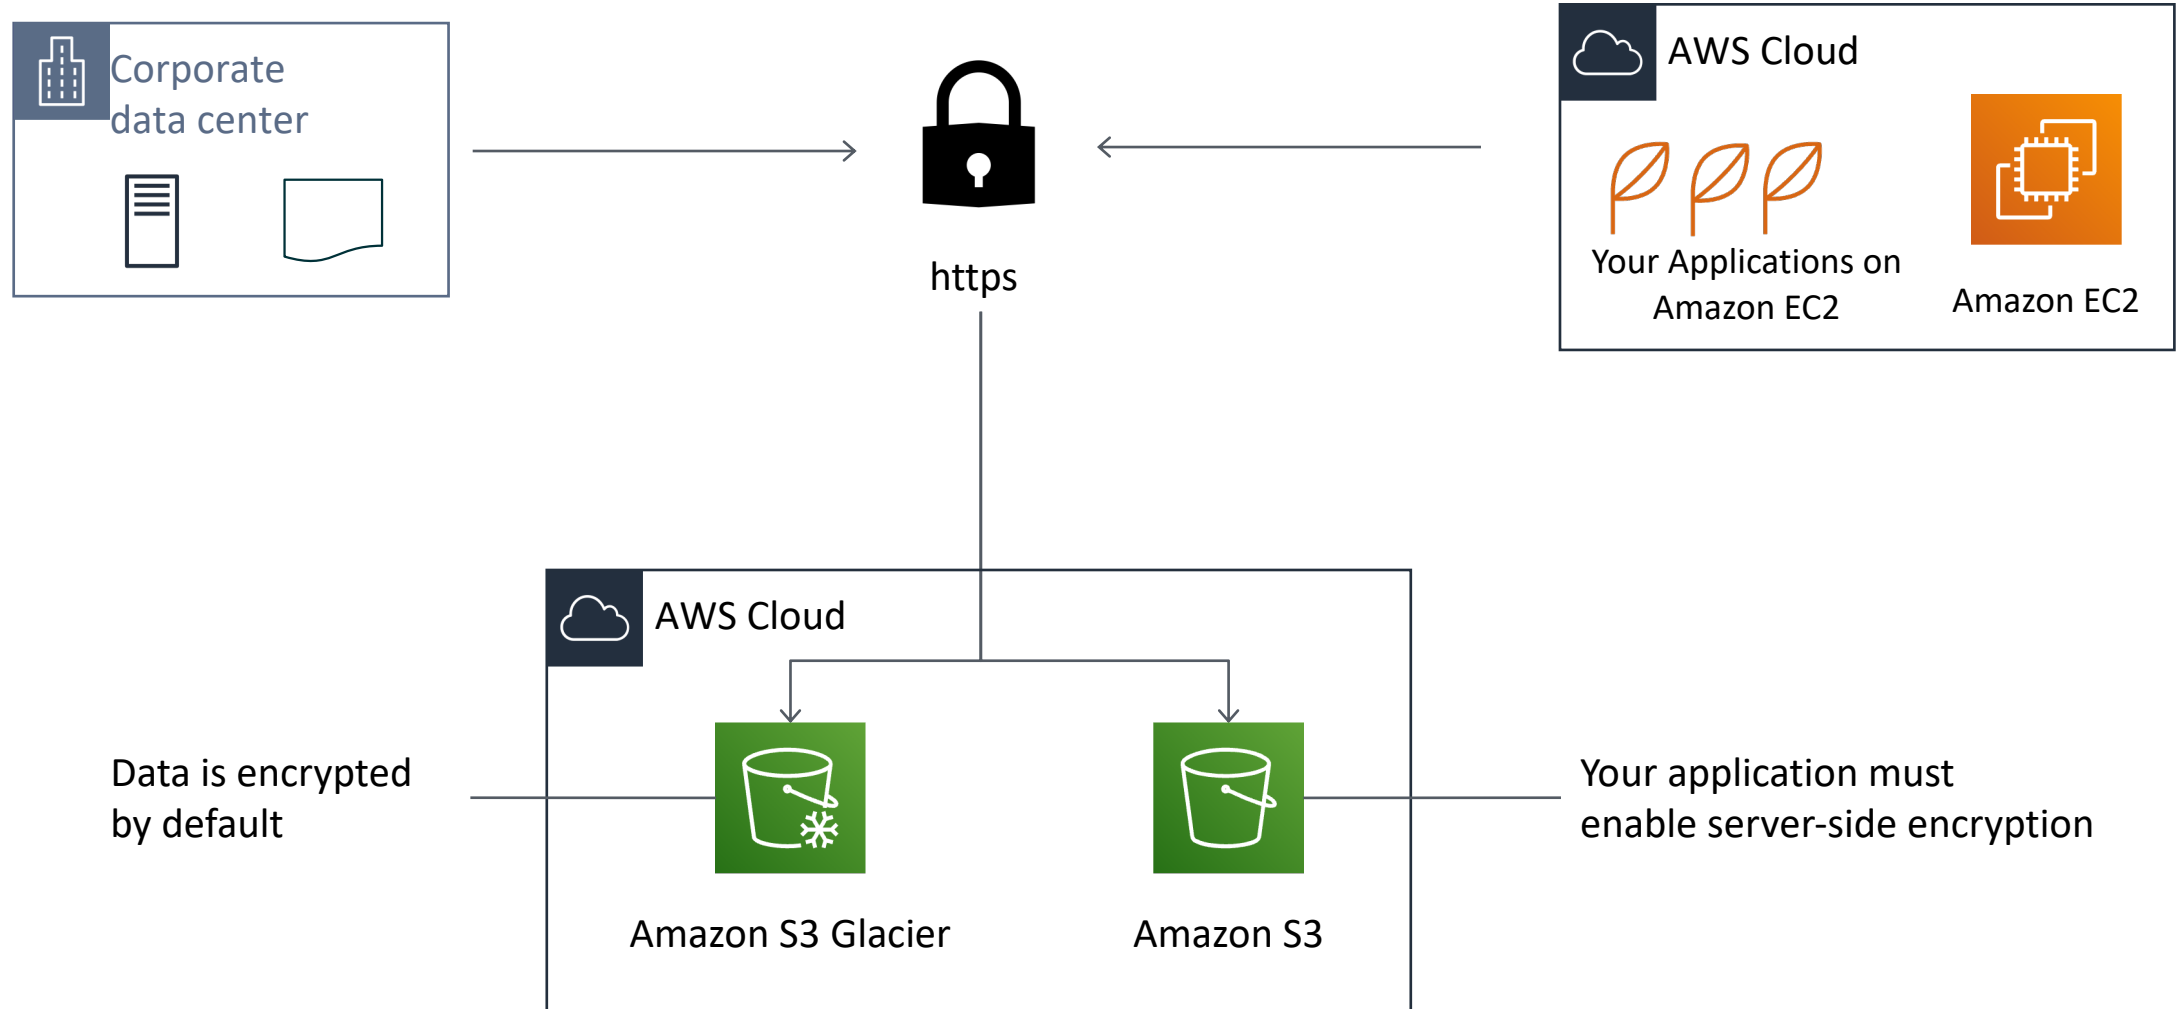
Amazon S3 lifecycle policies enable you to delete or move objects based on age.



Storage comparison

	Amazon S3	Amazon S3 Glacier
Data Volume	No limit	No limit
Average Latency	ms	minutes/hours
Item Size	5 TB maximum	40 TB maximum
Cost/GB per Month	Higher cost	Lower cost
Billed Requests	PUT, COPY, POST, LIST, and GET	UPLOAD and retrieval
Retrieval Pricing	¢ Per request	¢¢ Per request and per GB

Server-side encryption



Security with Amazon S3 Glacier



**Amazon S3
Glacier**



**Control access with
IAM**



**Amazon S3 Glacier encrypts
your data with AES-256**



**Amazon S3 Glacier manages
your keys for you**

Thank you

All trademarks are the property of their owners.

