K-Nearest Neighbors

Original Slides by:
Courtney Anne Ngo
Daniel Stanley Tan, PhD
Arren Antioquia

Updated (AY 2023 – 2024 T3) by:
Thomas James Tiam-Lee, PhD

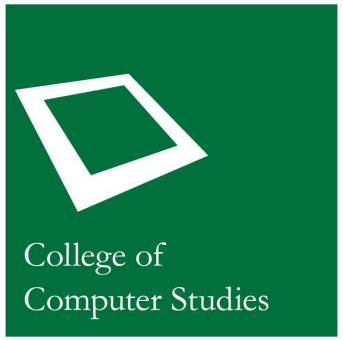

K-Nearest Neighbors (KNN)

- The most “naïve” kind of **supervised** machine learning model.
- It makes a prediction based on similarity to its closest neighbors.

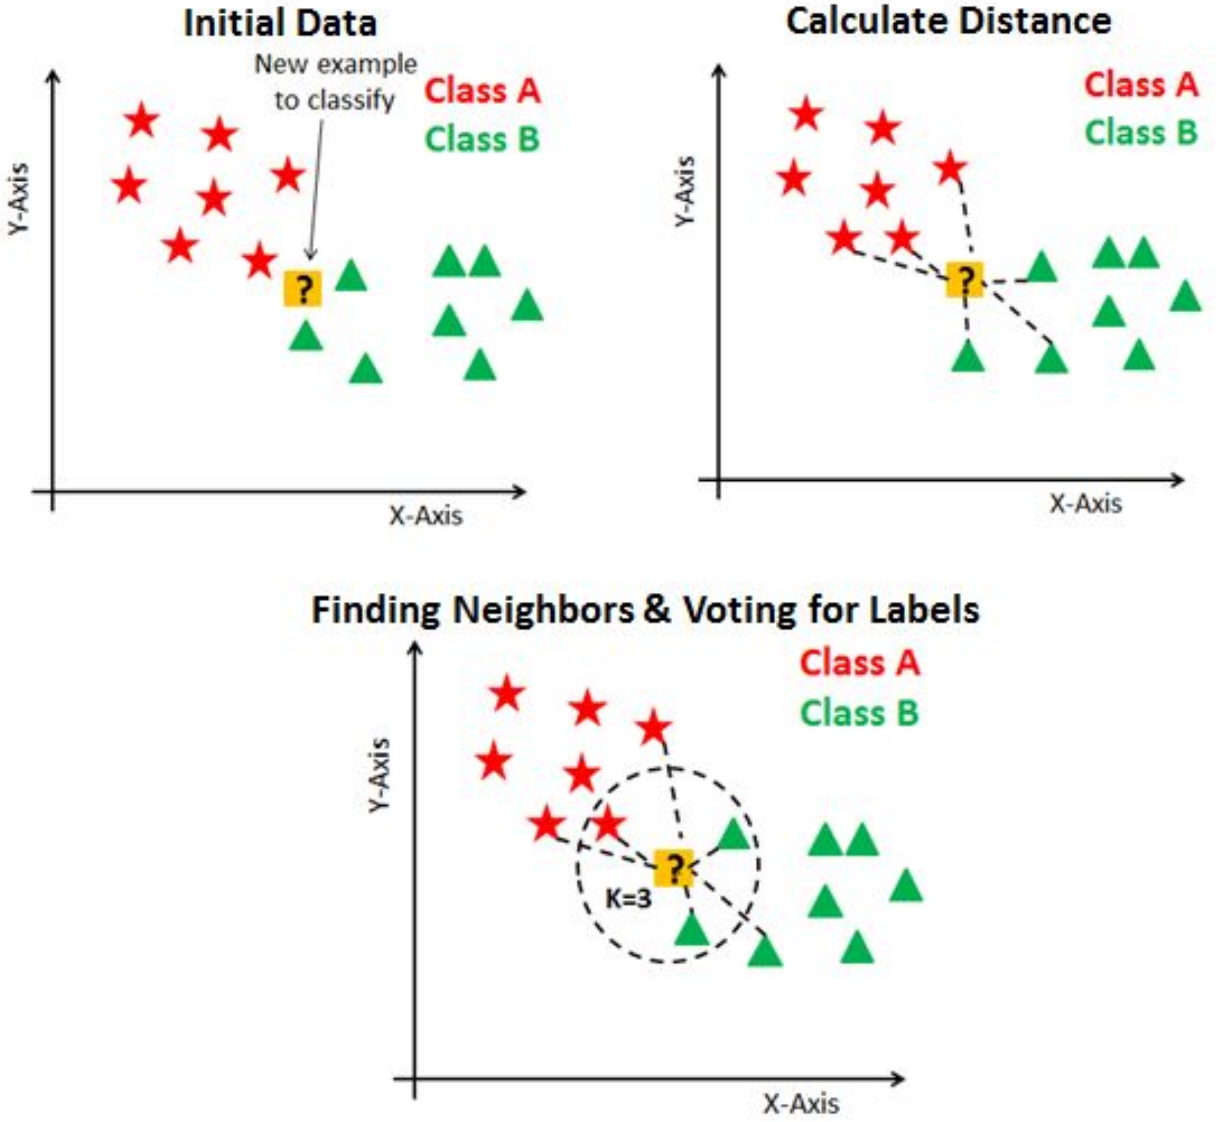

Sample Data 1

X y

temperature	humidity	weather
1	24	snowy
8	30	snowy
7	21	snowy
22	30	snowy
5	14	sunny
20	10	sunny
16	4	sunny
26	23	rainy
21	25	rainy
17	14	rainy
34	29	rainy

KNN: Training

- Not much “training” to be done.
- KNN simply **memorizes the entire dataset** and uses that as the model!

temperature	humidity	weather
1	24	snowy
8	30	snowy
7	21	snowy
22	30	snowy
5	14	sunny
20	10	sunny
16	4	sunny
26	23	rainy
21	25	rainy
17	14	rainy
34	29	rainy

Sample Data 1

x y

temperature	humidity	weather
1	24	snowy
8	30	snowy
7	21	snowy
22	30	snowy
5	14	sunny
20	10	sunny
16	4	sunny
26	23	rainy
21	25	rainy
17	14	rainy
34	29	rainy

temp	humidity	weather
21	17	?

temp	humidity	weather
1	24	snowy
8	30	snowy
7	21	snowy
22	30	snowy
5	14	sunny
20	10	sunny
16	4	sunny
26	223	rainy
21	25	rainy
17	14	rainy
34	29	rainy

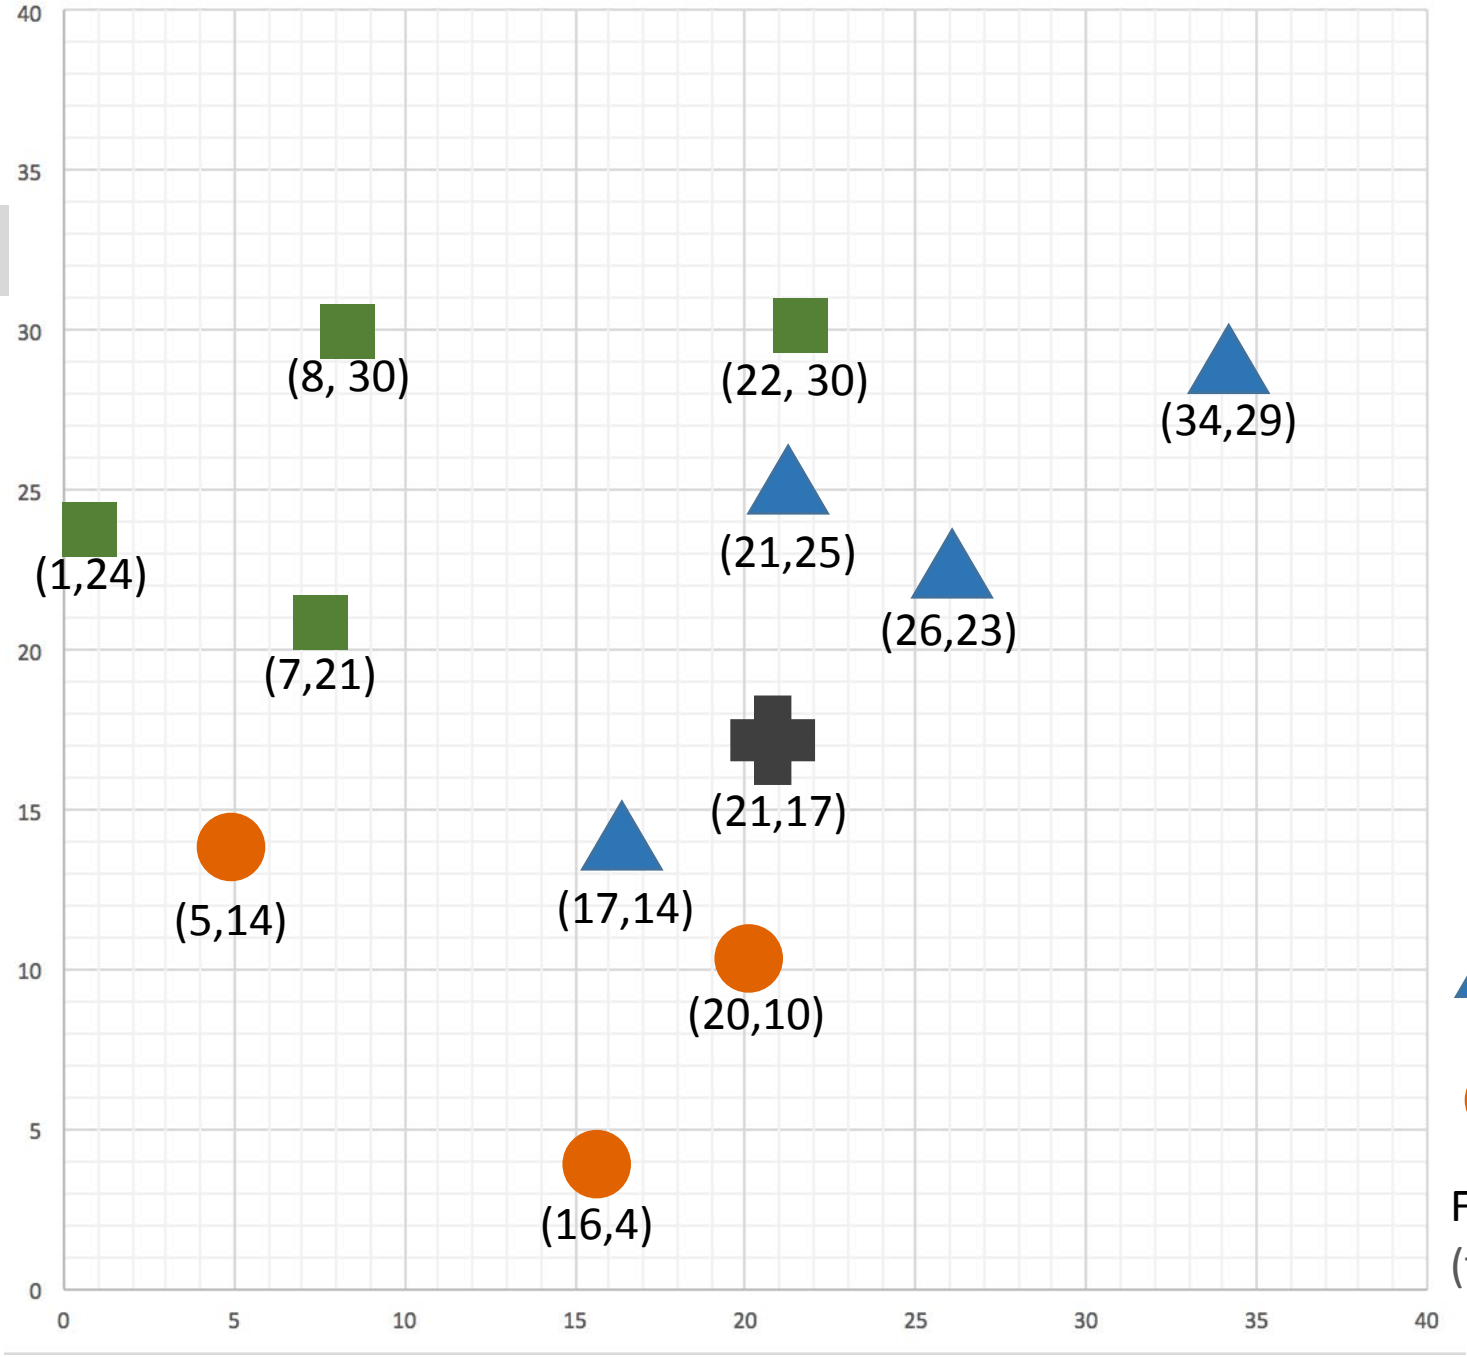

KNN: Prediction

- To make a prediction on an unknown instance, **find the most similar object and copy its class label!**

temperature	humidity	weather
1	24	snowy
8	30	snowy
7	21	snowy
22	30	snowy
5	14	sunny
20	10	sunny
16	4	sunny
26	23	rainy
21	25	rainy
17	14	rainy
34	29	rainy

temp	humidity	weather
21	17	?

Measures of Similarity

- Similarity of an instance z to the i-th training instance $X^{(i)}$
- **Euclidean Distance (L2-Distance)**:

$$dist\left(z, X^{(i)}\right) = \sqrt{\sum_{j=1}^{d} \left(z_j - X_j^{(i)}\right)^2}$$

- **Manhattan Distance (L1-Distance)**:

$$dist\left(z, X^{(i)}\right) = \sum_{j=1}^{d} \left|z_j - X_j^{(i)}\right|$$

temp	humidity	weather
1	24	snowy
8	30	snowy
7	21	snowy
22	30	snowy
5	14	sunny
20	10	sunny
16	4	sunny
26	223	rainy
21	25	rainy
17	14	rainy
34	29	rainy

Distances of each training data to ✚ (21,17)

(1,24)	(8, 30)	(7,21)	(22, 30)	(5,14)	(20,10)	(16,4)	(26,23)	(21,25)	(17,14)	(34,29)
21.19	18.38	14.56	13.03	16.28	7.07	13.93	7.81	8.00	5.00	17.69

$k = 1$

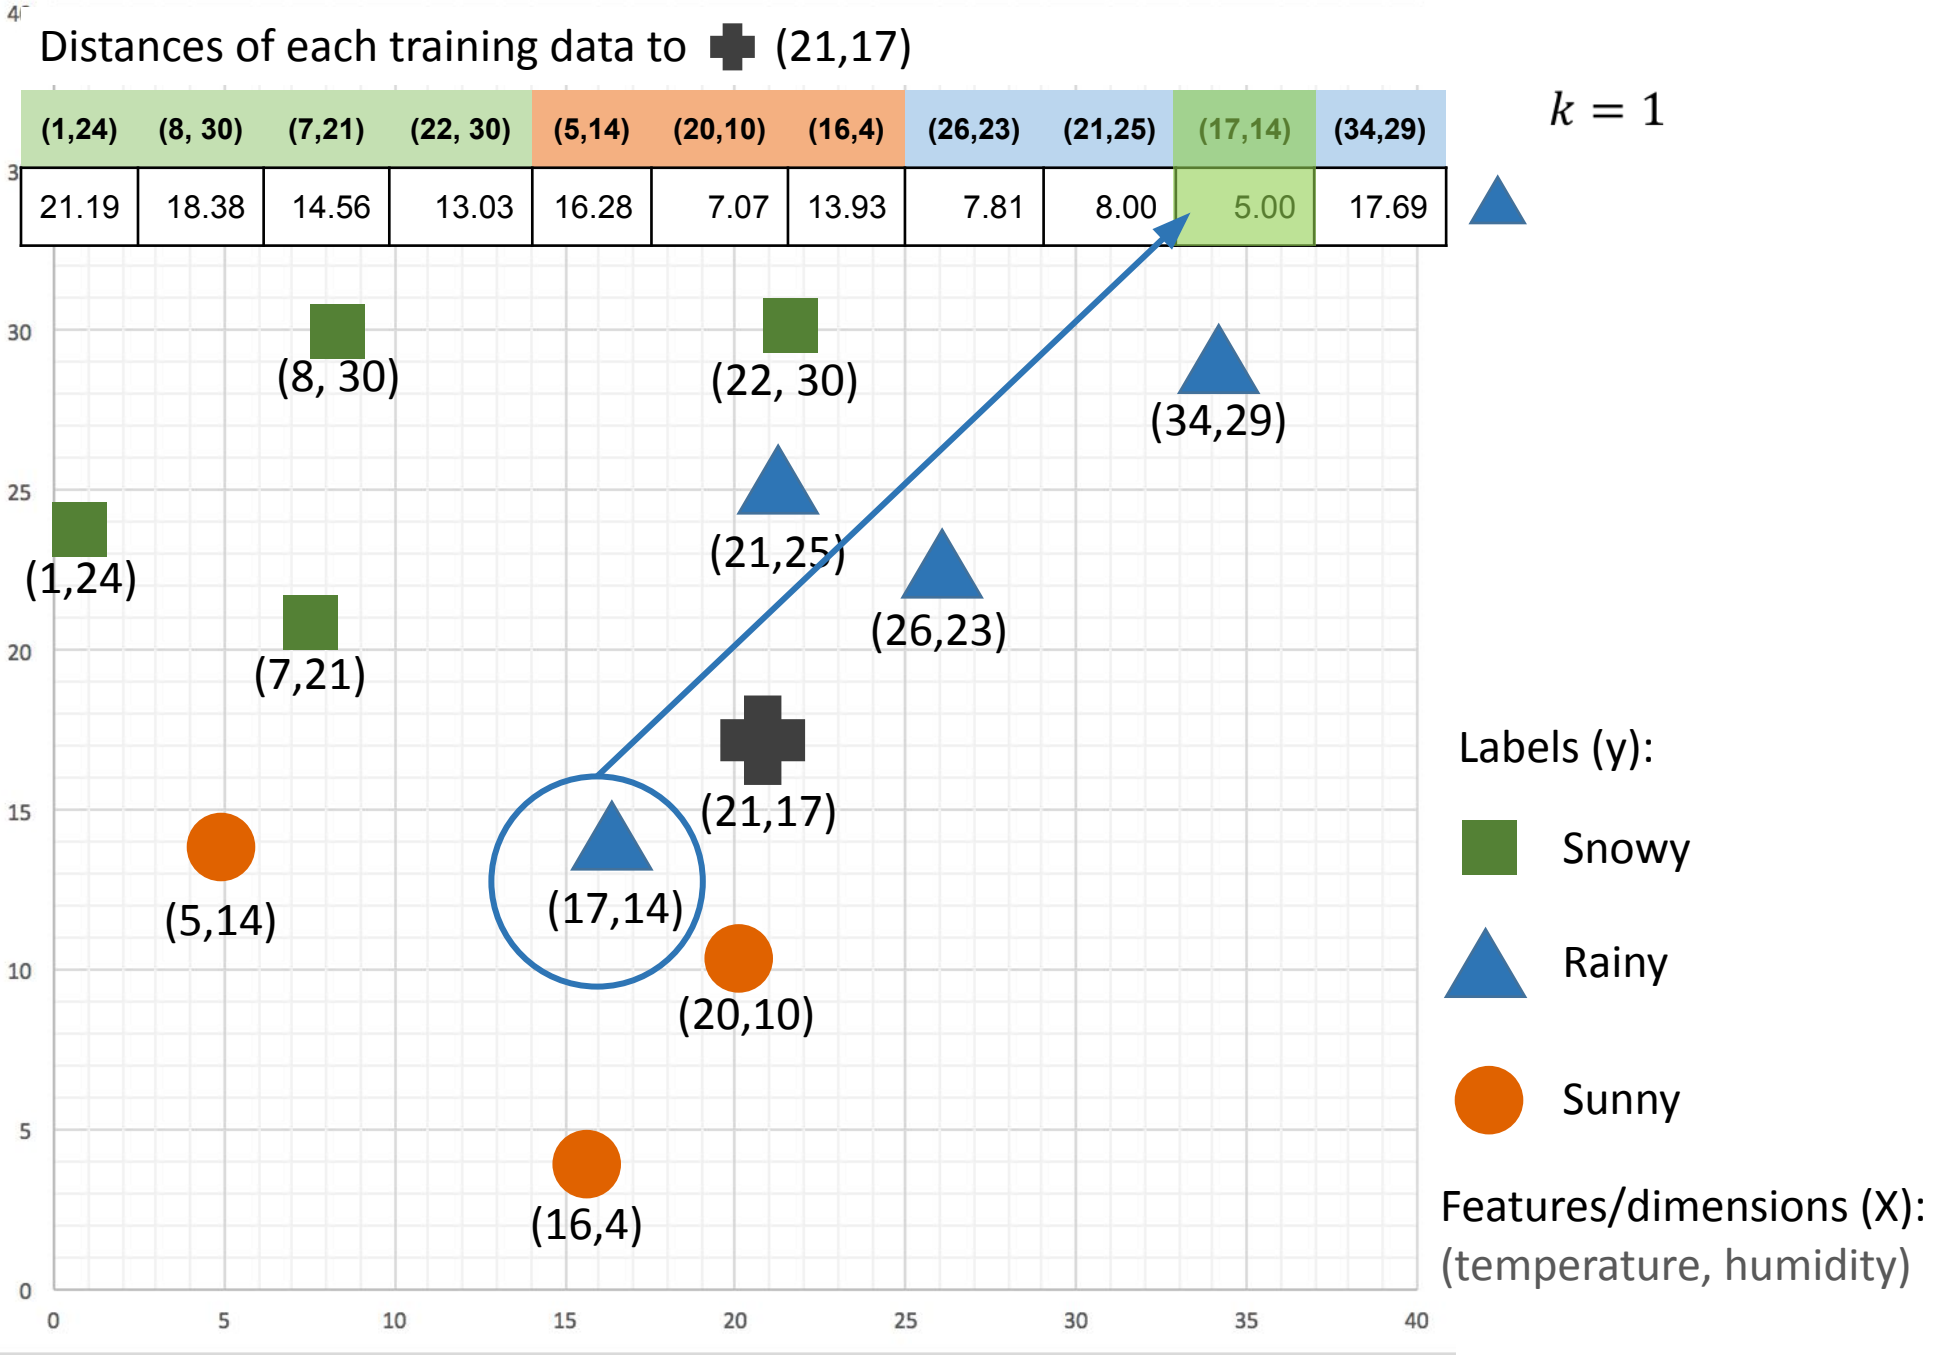

Labels (y):

Snowy

Rainy

Sunny

Features/dimensions (X):
(temperature, humidity)

The Hyperparameter k

- A **hyperparameter** is an option that you **manually decide on** when training an ML model.

- $\boldsymbol{k}$ is the **number of nearest neighbors** to consider before making a prediction.

temp	humidity	weather
1	24	snowy
8	30	snowy
7	21	snowy
22	30	snowy
5	14	sunny
20	10	sunny
16	4	sunny
26	223	rainy
21	25	rainy
17	14	rainy
34	29	rainy

Distances of each training data to (21,17)

	(1,24)	(8, 30)	(7,21)	(22, 30)	(5,14)	(20,10)	(16,4)	(26,23)	(21,25)	(17,14)	(34,29)
(21,17)	21.19	18.38	14.56	13.03	16.28	7.07	13.93	7.81	8.00	5.00	17.69
(33,13)	33.84	30.23	27.20	20.24	28.01	13.34	19.23	12.20	16.97	16.03	16.03
(7,6)	18.97	24.02	15.00	28.30	8.24	13.60	9.21	25.49	23.60	12.80	35.46

$k = 3$

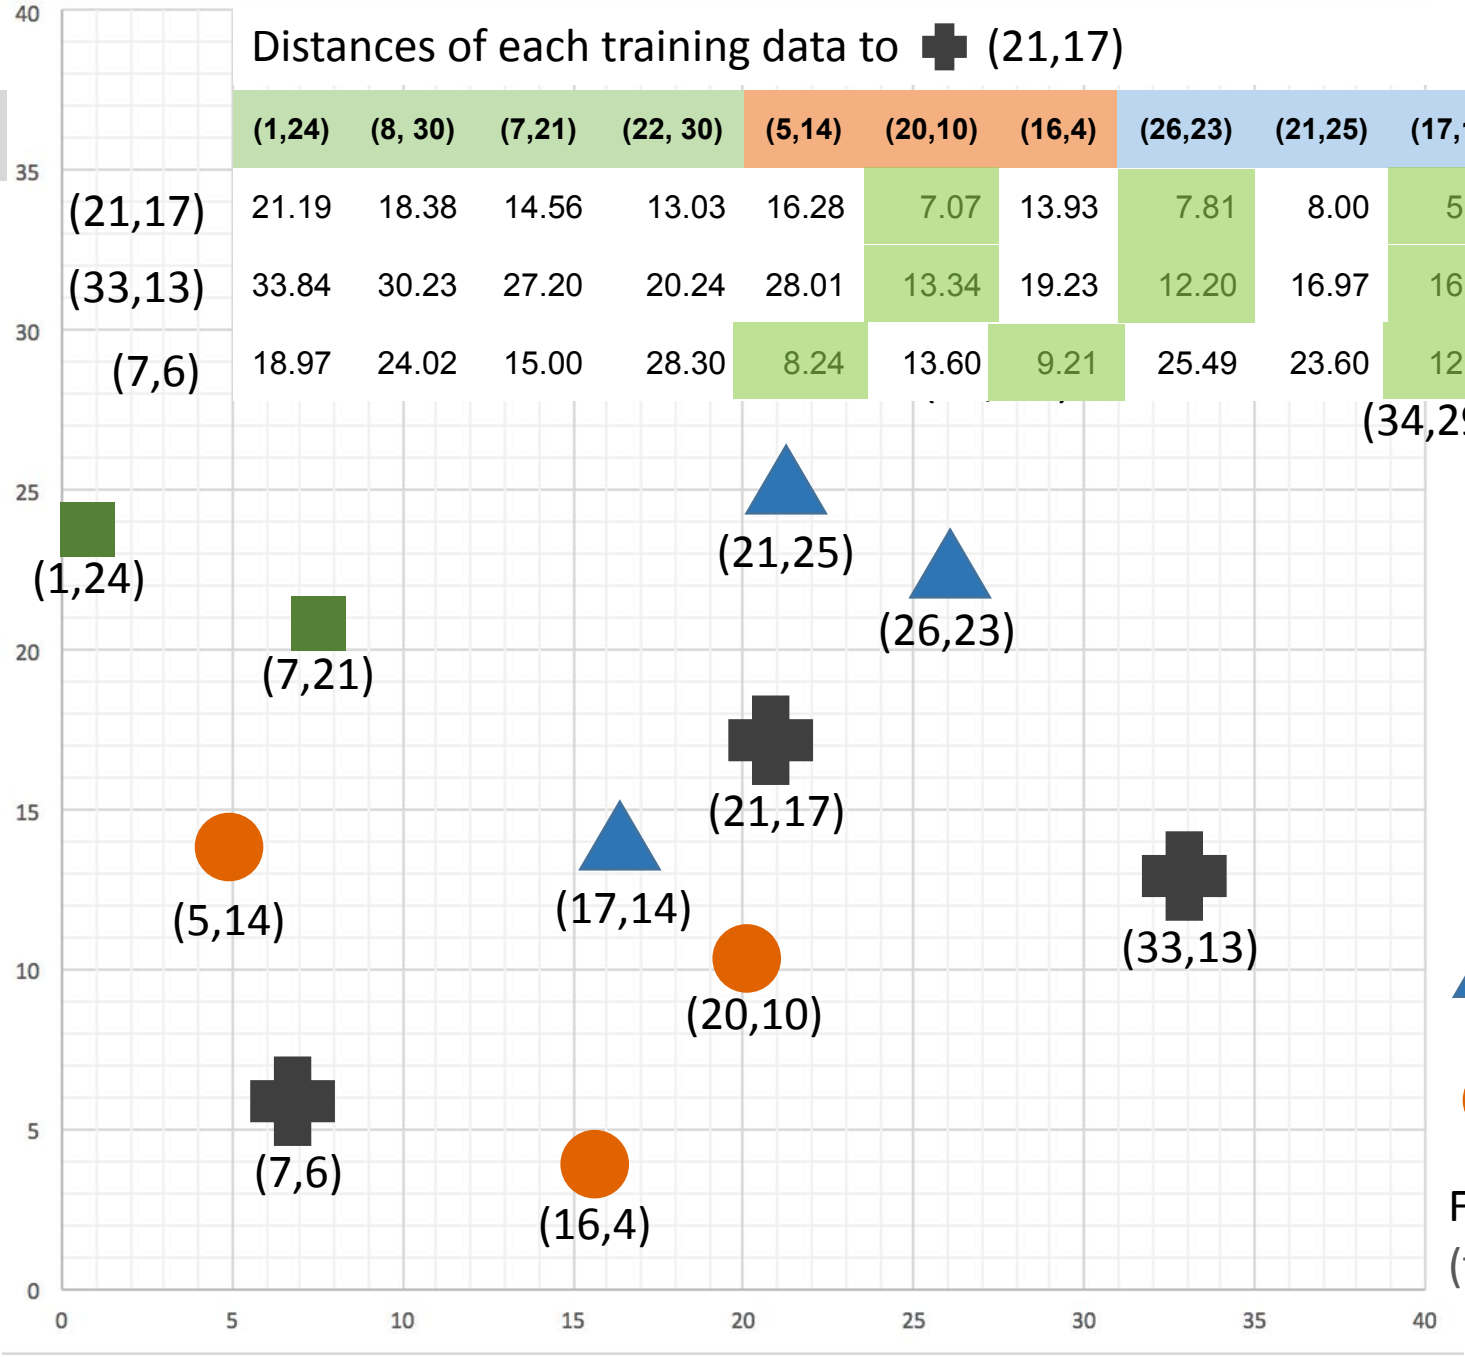

Labels (y):

Snowy

Rainy

Sunny

Features/dimensions (X):
(temperature, humidity)

Sample Data 2

x y

temperature	humidity	bacterial growth
1	24	4
8	30	10
7	21	6
22	30	10
5	14	2
20	10	8
16	4	4
26	23	8
21	25	10
17	14	10
34	29	2

temp	humidity	bacterial growth
21	17	?

Deciding the Final Prediction

- If $k > 1$…

For **classification**, majority wins! (mode)

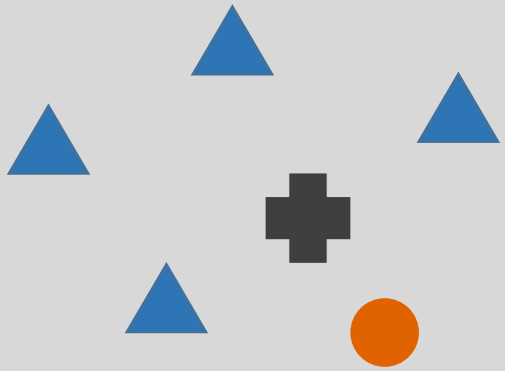

Final prediction: ▲

For **regression**, ????????

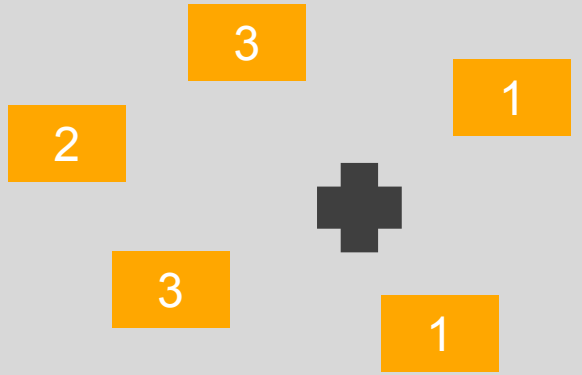

Final classification: ?

Deciding the Final Prediction

- If $k > 1$…

For **classification**, majority wins! (mode)

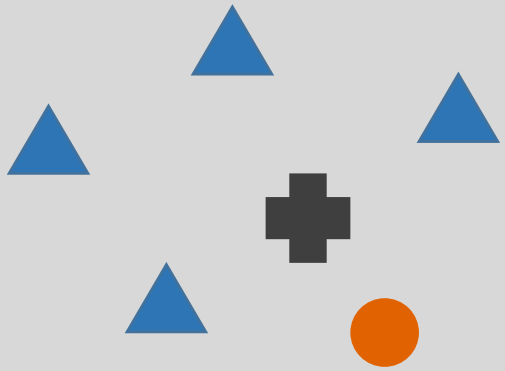

Final prediction: ▲

For **regression**, get the average!

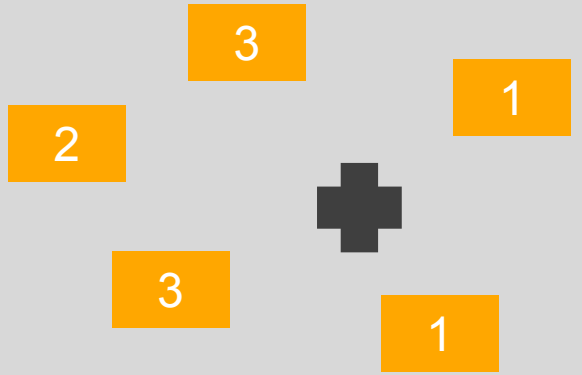

Final classification: 2

tempera ture	humidit y	bacterial growth
1	24	4
8	30	10
7	21	6
22	30	10
5	14	2
20	10	8
16	4	4
26	23	8
21	25	10
17	14	10
34	29	2

Distances of each training data to the test data

$k = 3$

	(1,24)	(8, 30)	(7,21)	(22, 30)	(5,14)	(20,10)	(16,4)	(26,23)	(21,25)	(17,14)	(34,29)	
(21,17)	21.19	18.38	14.56	13.03	16.28	7.07	13.93	7.81	8.00	5.00	17.69	8.667
(33,13)	33.84	30.23	27.20	20.24	28.01	13.34	19.23	12.20	16.97	16.03	16.03	8.667
(7,6)	18.97	24.02	15.00	28.30	8.24	13.60	9.21	25.49	23.60	12.80	35.46	5.667

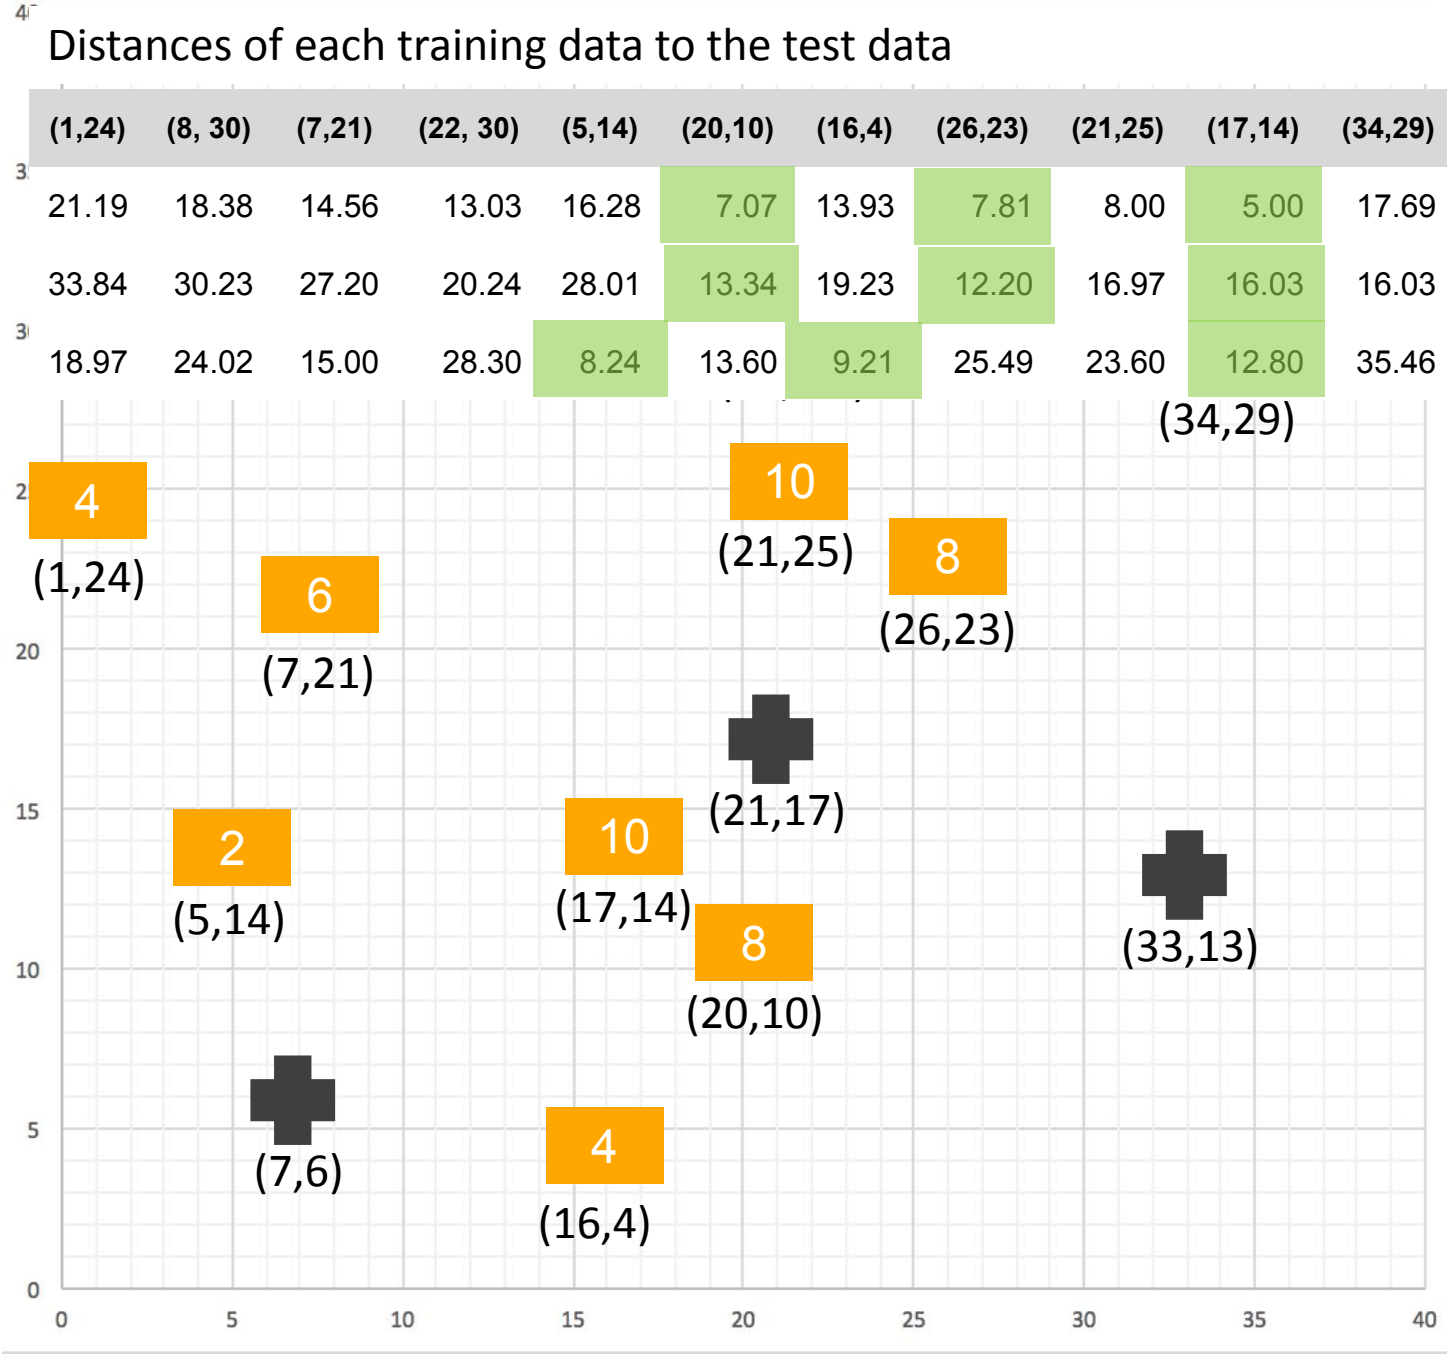

Features/dimensions (X):
(temperature, humidity)

K-Nearest Neighbors (KNN)

- Assumes all dimensions correspond to points in a d-dimensional space $\mathbb{R}^d$.
  - temperature, humidity ($\mathbb{R}^2$)

- **Features** may be discrete or continuous.

- **Labels** can be continuous (regression) or categorical (classification) as well.

The Distance Function Matters!

- Other less-commonly used distance functions:

  - **Minkowski Distance**
    - generalization of Euclidean and Manhattan distance
  - **Cosine distance**
    - similarity between two vectors
  - **Hamming distance**
    - similarity between two strings

How to Choose k?

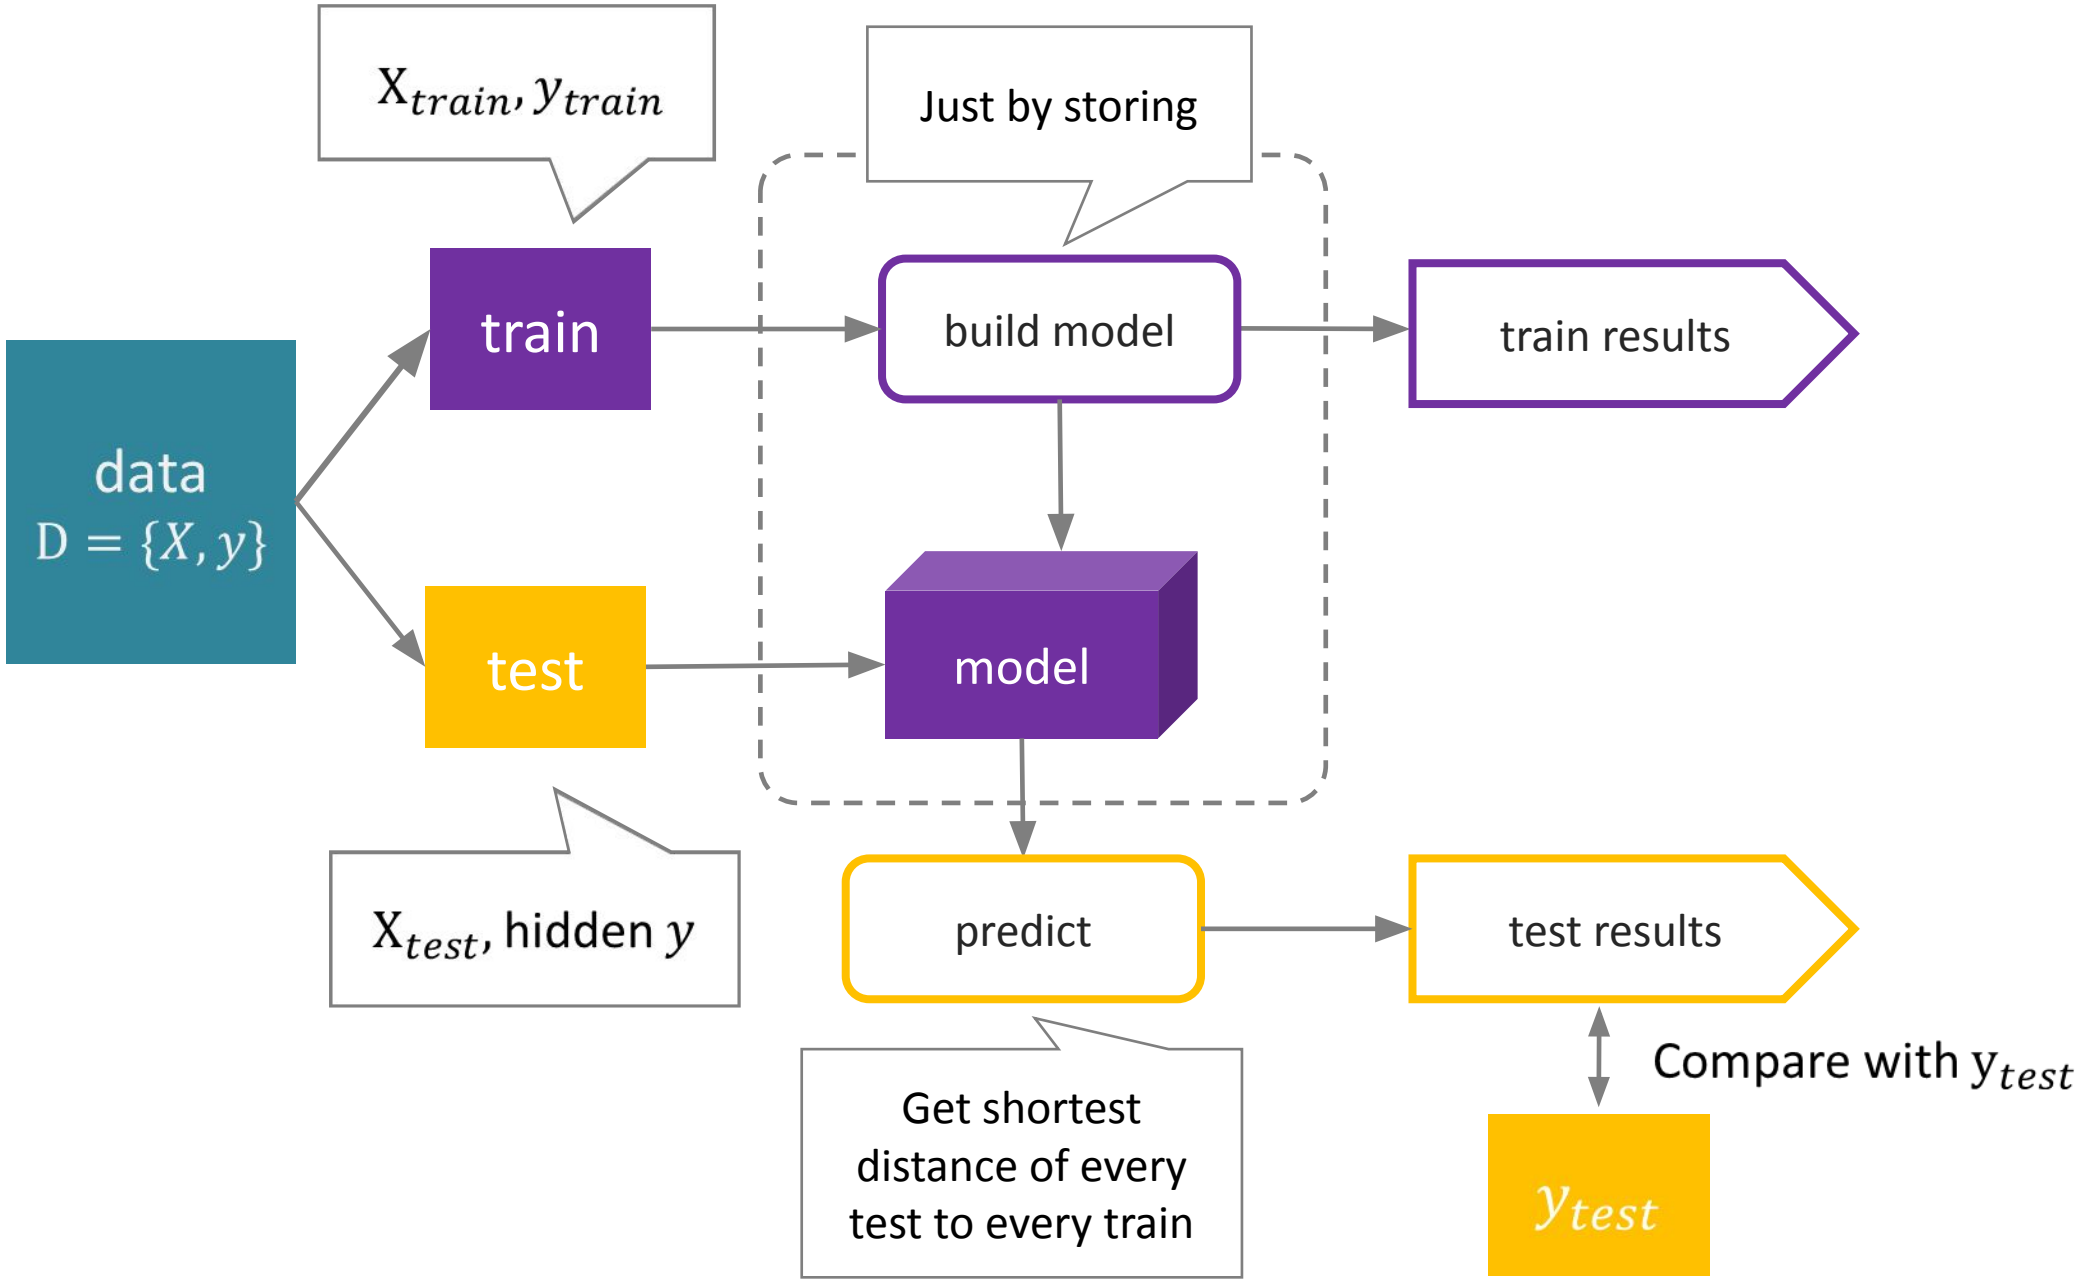
X_{train}, y_{train}
Just by storing
train
build model
train results
data
$D = \{X, y\}$
test
model
X_{test}, hidden y
predict
test results
Compare with y_{test}
Get shortest distance of every test to every train
y_{test}

Hyperparameter Tuning

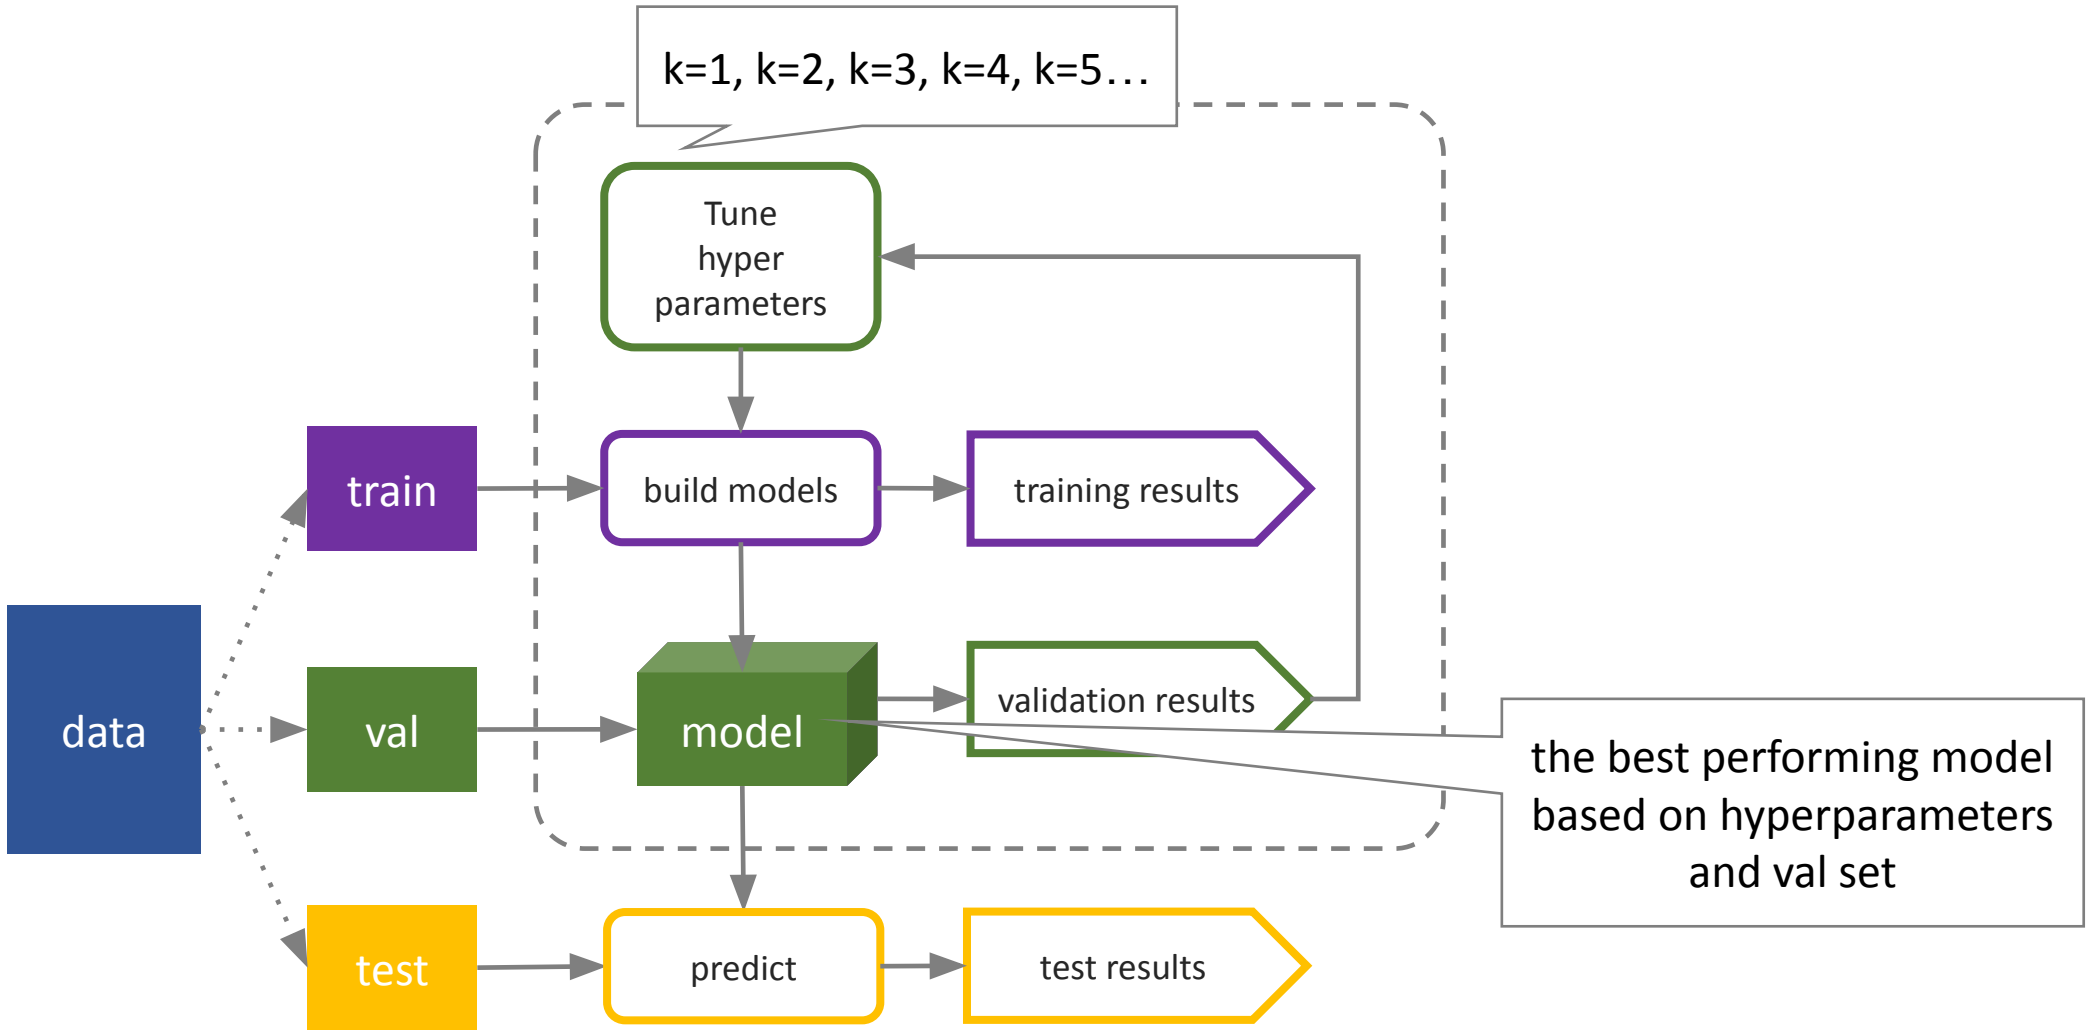

Hyperparameter Tuning with Cross-fold Validation

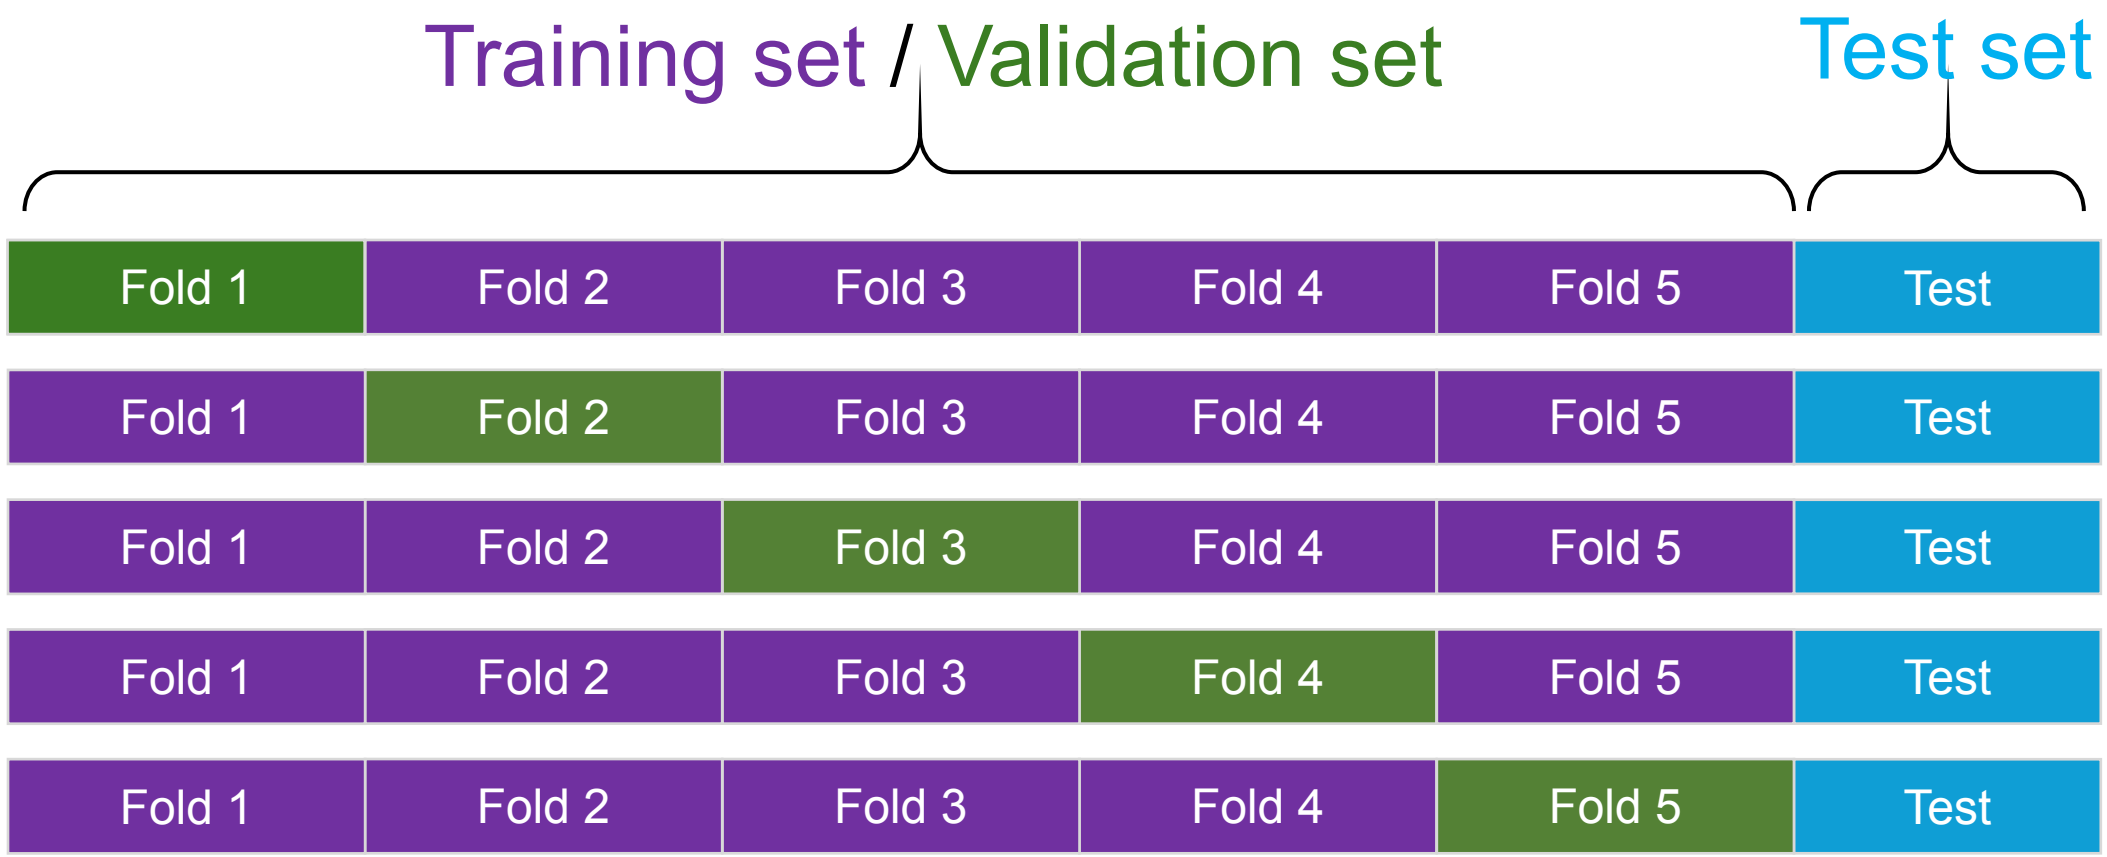

KNN Advantages and Disadvantages

- **Advantages**
  - Fast training time
  - Straightforward and easy to implement
- **Disadvantages**
  - Model is large
  - Prediction is slow if dataset is large
  - Considers all features equally, regardless of whether they are relevant or not